

Assignment 2: Implementing Feedforward neural networks with Keras and TensorFlow

```
In [ ]: #Tushar Santosh Kokane  
#B511066
```

```
In [2]: #installations  
from sklearn.preprocessing import LabelBinarizer  
from sklearn.metrics import classification_report  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense  
from tensorflow.keras.optimizers import SGD  
from tensorflow.keras.datasets import mnist  
from tensorflow.keras import backend as K  
import matplotlib.pyplot as plt  
import numpy as np
```

```
In [2]: #grabbing the mnist dataset  
((X_train, Y_train), (X_test, Y_test)) = mnist.load_data()  
X_train = X_train.reshape((X_train.shape[0], 28 * 28 * 1))  
X_test = X_test.reshape((X_test.shape[0], 28 * 28 * 1))  
X_train = X_train.astype("float32") / 255.0  
X_test = X_test.astype("float32") / 255.0
```

```
In [3]: lb = LabelBinarizer()  
Y_train = lb.fit_transform(Y_train)  
Y_test = lb.transform(Y_test)
```

```
In [4]: #building the model  
model = Sequential()  
model.add(Dense(128, input_shape=(784,), activation="sigmoid"))  
model.add(Dense(64, activation="sigmoid"))  
model.add(Dense(10, activation="softmax"))
```

```
In [5]: sgd = SGD(0.01)
epochs=10
model.compile(loss="categorical_crossentropy", optimizer=sgd,metrics=["accu
H = model.fit(X_train, Y_train, validation_data=(X_test, Y_test),epochs=epo
```

```
Epoch 1/10
469/469 [=====] - 9s 16ms/step - loss: 2.2897 - a
ccuracy: 0.1692 - val_loss: 2.2521 - val_accuracy: 0.2864
Epoch 2/10
469/469 [=====] - 6s 12ms/step - loss: 2.2256 - a
ccuracy: 0.3262 - val_loss: 2.1920 - val_accuracy: 0.4748
Epoch 3/10
469/469 [=====] - 5s 11ms/step - loss: 2.1567 - a
ccuracy: 0.4901 - val_loss: 2.1095 - val_accuracy: 0.5230
Epoch 4/10
469/469 [=====] - 5s 11ms/step - loss: 2.0593 - a
ccuracy: 0.5760 - val_loss: 1.9911 - val_accuracy: 0.6251
Epoch 5/10
469/469 [=====] - 5s 12ms/step - loss: 1.9214 - a
ccuracy: 0.6266 - val_loss: 1.8291 - val_accuracy: 0.6580
Epoch 6/10
469/469 [=====] - 5s 11ms/step - loss: 1.7453 - a
ccuracy: 0.6672 - val_loss: 1.6372 - val_accuracy: 0.6906
Epoch 7/10
469/469 [=====] - 5s 11ms/step - loss: 1.5541 - a
ccuracy: 0.6961 - val_loss: 1.4464 - val_accuracy: 0.7194
Epoch 8/10
469/469 [=====] - 5s 12ms/step - loss: 1.3752 - a
ccuracy: 0.7232 - val_loss: 1.2776 - val_accuracy: 0.7429
Epoch 9/10
469/469 [=====] - 5s 12ms/step - loss: 1.2225 - a
ccuracy: 0.7451 - val_loss: 1.1384 - val_accuracy: 0.7587
Epoch 10/10
469/469 [=====] - 5s 11ms/step - loss: 1.0970 - a
ccuracy: 0.7626 - val_loss: 1.0251 - val_accuracy: 0.7720
```

```
In [6]: #making the predictions
predictions = model.predict(X_test, batch_size=128)
print(classification_report(Y_test.argmax(axis=1),predictions.argmax(axis=1)
```

```
79/79 [=====] - 0s 3ms/step
              precision    recall  f1-score   support

     0           0.83         0.96         0.89         980
     1           0.76         0.99         0.86        1135
     2           0.82         0.64         0.72        1032
     3           0.65         0.85         0.74        1010
     4           0.71         0.83         0.76         982
     5           0.83         0.38         0.52         892
     6           0.85         0.89         0.87         958
     7           0.82         0.88         0.85        1028
     8           0.79         0.63         0.70         974
     9           0.75         0.61         0.67        1009

 accuracy          0.77         0.77         0.77        10000
 macro avg          0.78         0.77         0.76        10000
 weighted avg       0.78         0.77         0.76        10000
```

```
In [8]: #plotting the training loss and accuracy
plt.style.use("ggplot")
plt.figure()
plt.plot(np.arange(0, epochs), H.history["loss"], label="train_loss")
plt.plot(np.arange(0, epochs), H.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, epochs), H.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, epochs), H.history["val_accuracy"], label="val_acc")
plt.title("Training Loss and Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy")
plt.legend()
```

Out[8]: <matplotlib.legend.Legend at 0x245819a3e50>

