

Assignment 4: ECG Anomaly detection using Autoencoders

In []: *#Tushar Kokane*
#B511066 Div: A


In [2]: *#importing libraries and dataset*
`import numpy as np`
`import pandas as pd`
`import tensorflow as tf`
`import matplotlib.pyplot as plt`
`from sklearn.metrics import accuracy_score`
`from tensorflow.keras.optimizers import Adam`
`from sklearn.preprocessing import MinMaxScaler`
`from tensorflow.keras import Model, Sequential`
`from tensorflow.keras.layers import Dense, Dropout`
`from sklearn.model_selection import train_test_split`
`from tensorflow.keras.losses import MeanSquaredLogarithmicError`

`PATH_TO_DATA = 'http://storage.googleapis.com/download.tensorflow.org/data/'`
`data = pd.read_csv(PATH_TO_DATA, header=None)`
`data.head()`

Out[2]:

	0	1	2	3	4	5	6	7	
0	-0.112522	-2.827204	-3.773897	-4.349751	-4.376041	-3.474986	-2.181408	-1.818286	-1.25
1	-1.100878	-3.996840	-4.285843	-4.506579	-4.022377	-3.234368	-1.566126	-0.992258	-0.75
2	-0.567088	-2.593450	-3.874230	-4.584095	-4.187449	-3.151462	-1.742940	-1.490659	-1.18
3	0.490473	-1.914407	-3.616364	-4.318823	-4.268016	-3.881110	-2.993280	-1.671131	-1.33
4	0.800232	-0.874252	-2.384761	-3.973292	-4.338224	-3.802422	-2.534510	-1.783423	-1.59

5 rows × 141 columns



In [2]: *#finding shape of the dataset*
`data.shape`

Out[2]: (4998, 141)

In [8]: *#splitting training and testing dataset*
`features = data.drop(140, axis=1)`
`target = data[140]`
`x_train, x_test, y_train, y_test = train_test_split(`
 `features, target, test_size=0.2, stratify=target`
`)`
`train_index = y_train[y_train == 1].index`
`train_data = x_train.loc[train_index]`

```
In [10]: #scaling the data using MinMaxScaler
min_max_scaler = MinMaxScaler(feature_range=(0, 1))
x_train_scaled = min_max_scaler.fit_transform(train_data.copy())
x_test_scaled = min_max_scaler.transform(x_test.copy())
```

```
In [13]: #creating autoencoder subclass by extending Model class from keras
class AutoEncoder(Model):
    def __init__(self, output_units, ldim=8):
        super().__init__()
        self.encoder = Sequential([
            Dense(64, activation='relu'),
            Dropout(0.1),
            Dense(32, activation='relu'),
            Dropout(0.1),
            Dense(16, activation='relu'),
            Dropout(0.1),
            Dense(ldim, activation='relu')
        ])
        self.decoder = Sequential([
            Dense(16, activation='relu'),
            Dropout(0.1),
            Dense(32, activation='relu'),
            Dropout(0.1),
            Dense(64, activation='relu'),
            Dropout(0.1),
            Dense(output_units, activation='sigmoid')
        ])


    def call(self, inputs):
        encoded = self.encoder(inputs)
        decoded = self.decoder(encoded)
        return decoded
```

```
In [16]: #model configuration
model = AutoEncoder(output_units=x_train_scaled.shape[1])
model.compile(loss='mse', metrics=['mse'], optimizer='adam')
epochs = 20

history = model.fit(
    x_train_scaled,
    x_train_scaled,
    epochs=epochs,
    batch_size=512,
    validation_data=(x_test_scaled, x_test_scaled)
)
```

```
Epoch 1/20
5/5 [=====] - 5s 152ms/step - loss: 0.0113 - mse:
0.0248 - val_loss: 0.0132 - val_mse: 0.0303
Epoch 2/20
5/5 [=====] - 0s 38ms/step - loss: 0.0110 - mse:
0.0242 - val_loss: 0.0131 - val_mse: 0.0299
Epoch 3/20
5/5 [=====] - 0s 43ms/step - loss: 0.0104 - mse:
0.0228 - val_loss: 0.0128 - val_mse: 0.0294
Epoch 4/20
5/5 [=====] - 0s 45ms/step - loss: 0.0094 - mse:
0.0206 - val_loss: 0.0129 - val_mse: 0.0296
Epoch 5/20
5/5 [=====] - 0s 39ms/step - loss: 0.0083 - mse:
0.0184 - val_loss: 0.0126 - val_mse: 0.0289
Epoch 6/20
5/5 [=====] - 0s 43ms/step - loss: 0.0073 - mse:
0.0161 - val_loss: 0.0118 - val_mse: 0.0273
Epoch 7/20
5/5 [=====] - 0s 45ms/step - loss: 0.0065 - mse:
0.0143 - val_loss: 0.0114 - val_mse: 0.0263
Epoch 8/20
5/5 [=====] - 0s 39ms/step - loss: 0.0060 - mse:
0.0131 - val_loss: 0.0109 - val_mse: 0.0252
Epoch 9/20
5/5 [=====] - 0s 39ms/step - loss: 0.0056 - mse:
0.0123 - val_loss: 0.0105 - val_mse: 0.0242
Epoch 10/20
5/5 [=====] - 0s 42ms/step - loss: 0.0053 - mse:
0.0116 - val_loss: 0.0101 - val_mse: 0.0233
Epoch 11/20
5/5 [=====] - 0s 39ms/step - loss: 0.0051 - mse:
0.0112 - val_loss: 0.0099 - val_mse: 0.0229
Epoch 12/20
5/5 [=====] - 0s 38ms/step - loss: 0.0049 - mse:
0.0109 - val_loss: 0.0096 - val_mse: 0.0223
Epoch 13/20
5/5 [=====] - 0s 45ms/step - loss: 0.0048 - mse:
0.0106 - val_loss: 0.0096 - val_mse: 0.0223
Epoch 14/20
5/5 [=====] - 0s 45ms/step - loss: 0.0047 - mse:
0.0104 - val_loss: 0.0095 - val_mse: 0.0221
Epoch 15/20
5/5 [=====] - 0s 41ms/step - loss: 0.0047 - mse:
0.0103 - val_loss: 0.0095 - val_mse: 0.0220
Epoch 16/20
5/5 [=====] - 0s 41ms/step - loss: 0.0046 - mse:
0.0102 - val_loss: 0.0095 - val_mse: 0.0220
Epoch 17/20
5/5 [=====] - 0s 39ms/step - loss: 0.0046 - mse:
0.0101 - val_loss: 0.0095 - val_mse: 0.0220
Epoch 18/20
5/5 [=====] - 0s 41ms/step - loss: 0.0045 - mse:
0.0100 - val_loss: 0.0094 - val_mse: 0.0220
Epoch 19/20
5/5 [=====] - 0s 41ms/step - loss: 0.0045 - mse:
0.0100 - val_loss: 0.0094 - val_mse: 0.0219
Epoch 20/20
5/5 [=====] - 0s 40ms/step - loss: 0.0045 - mse:
0.0099 - val_loss: 0.0094 - val_mse: 0.0218
```

```
In [17]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.xlabel('Epochs')
plt.ylabel('MSLE Loss')
plt.legend(['loss', 'val_loss'])
plt.show()
```



```
In [18]: #finding threshold for anomaly and doing predictions
def find_threshold(model, x_train_scaled):
    reconstructions = model.predict(x_train_scaled)
    reconstruction_errors = tf.keras.losses.msle(reconstructions, x_train_scaled)
    threshold = np.mean(reconstruction_errors.numpy()) \
        + np.std(reconstruction_errors.numpy())
    return threshold

def get_predictions(model, x_test_scaled, threshold):
    predictions = model.predict(x_test_scaled)
    errors = tf.keras.losses.msle(predictions, x_test_scaled)
    anomaly_mask = pd.Series(errors) > threshold
    preds = anomaly_mask.map(lambda x: 0.0 if x == True else 1.0)
    return preds

threshold = find_threshold(model, x_train_scaled)
print(f"Threshold: {threshold}")
```

```
73/73 [=====] - 1s 4ms/step
Threshold: 0.0099591106762396
```

```
In [19]: #getting accuracy score  
predictions = get_predictions(model, x_test_scaled, threshold)  
accuracy_score(predictions, y_test)
```

32/32 [=====] - 0s 5ms/step

Out[19]: 0.935